

FINAL PROJECT:

**Hardware
Accelerator For
Cache Timing Attacks
Mitigations**



**Lidor Azani
Noam Hever**

Supervision: Dr. Itamar Levi, Mr. Oren Ganon



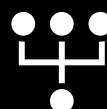
18 minutes

What will we show?



1

Short introduction –
overview of previous
presentation.
The solution we chose.



2

Project's plan
(waterfall) .
its development
and analysis
stages.
tools used.



3

The built solution:
Our Memory access
flow – in software
and hardware
accelerated.



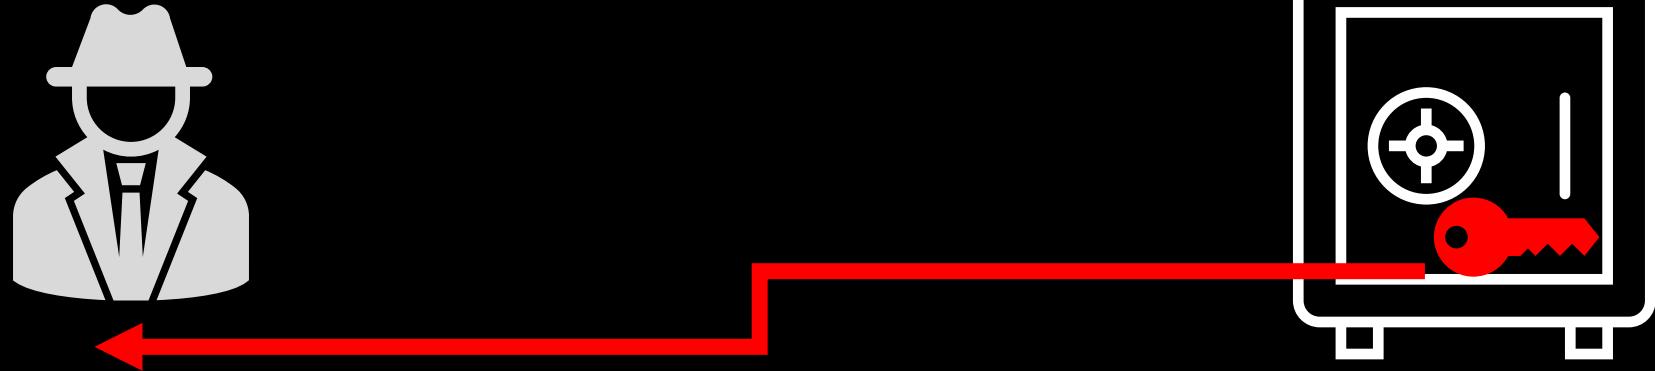
4

Analysis:

- a. Cryptographic – statistical analysis.
- b. Performance and acceleration analysis.

Introduction

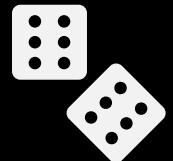
- During previous presentation we talked about:
 1. Side channels attacks .
 2. Cache timing attacks.
 3. Scramble cache – a defense mechanism.



Information about the secret is leaking
through a nonconventional information
channel/signal.

Introduction

- Attacker can learn information about cryptographic secrets by observing access time patterns to the cache/memory on specific sets.
- Different Mitigations are focusing on preventing that by adding randomness to the memory/cache accesses .



Introduction – A Reminder:

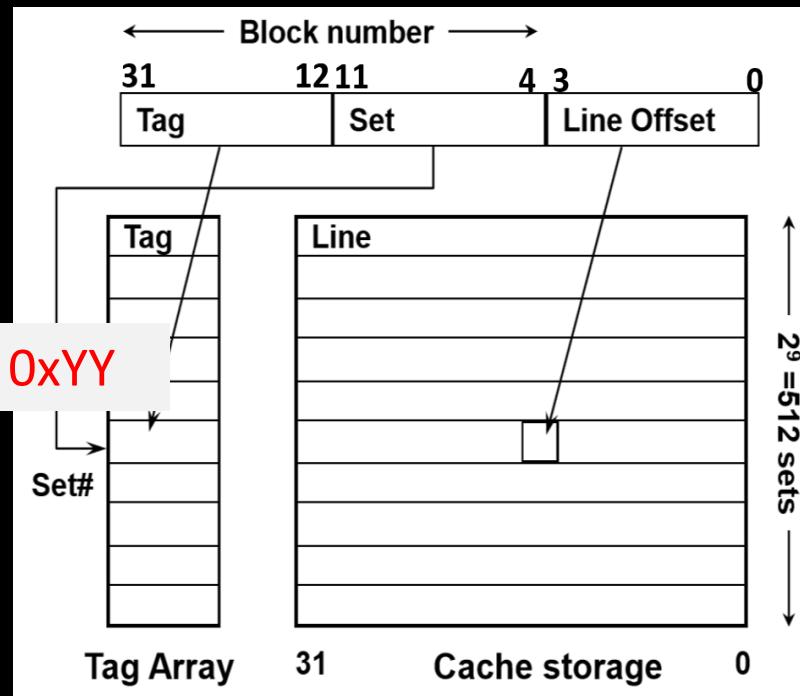
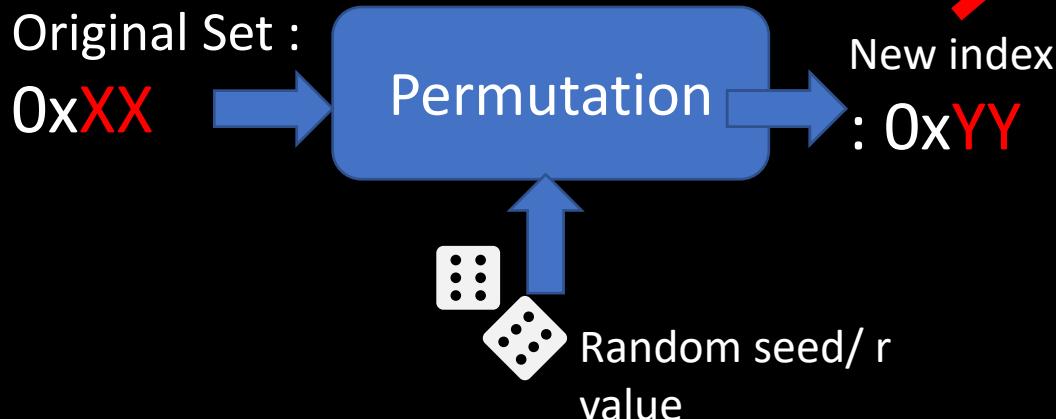
Random permutation cache – third solution

What can we do?

1. give each process a permutation table

Instead of using the “set” to decide on where we should bring the new data – we will decide using a permutation function on the value of the “set” in the address.

How do we calculate the mapping:



Introduction

Announced Goals:

1. Reduce the attacker's ability to learn anything about the memory accesses done by another process.
2. Randomize memory accesses in a cryptographically secure and efficient way .



2. Implementation

Inspired by the “Scrambled cache ” paper – with real implementation
and no abstraction



2. Implementation

The “Scrambled cache” paper :

1. - only abstract emulation - no implementation
2. - no “full picture” of cost
3. - no effectiveness demonstration – with simulating attacks



Our “Randomized Memory area solution”

1. - Software and hardware implementation –usable by anyone!
2. - Cost analysis with a synthesizable processor
3. - Effective demonstration – ability to any C / C++ program – either attacker/Victim simulation or benchmarks

2. Implementation: Development Stages:

Research papers

Planning software and hardware accelerations

C prototype



Validating: correctness, injectivity of permutation



In TIE HDL – Accelerate



Measure, test and benchmark

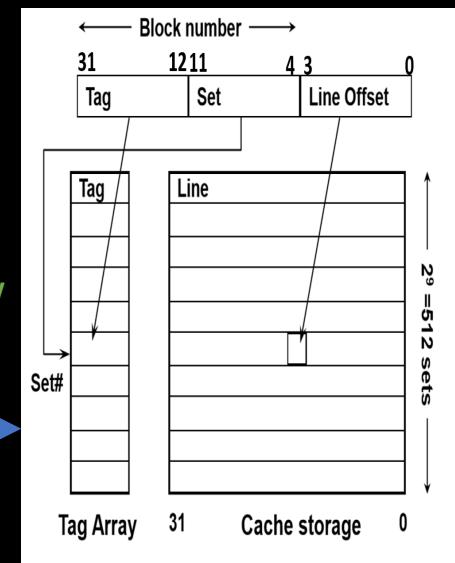


Prove functionality, safety and efficiency



Solution schematics

Regular L1 Cache

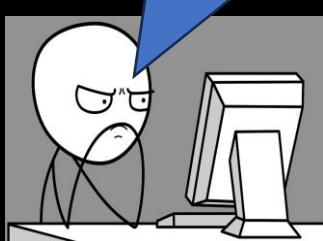


Permutation

New address

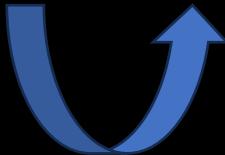
Original address

Read memory[10]



History table – compare sets

index	History table - original indexes
0	original set for memory line 0
1	original set for memory line 1
2	original set for memory line 2
3	original set for memory line 3
4	original set for memory line 4
5	original set for memory line 5



If sets are not matching: search history table for a match.
Then fix mapping by swapping lines in memory and update history table

Access memory with New address

On miss – access RAM

Memory area:

Memory arr:

```
line 000: 0x474F3000 : 0x00 0x00
line 001: 0x474F3010 : 0x00 0x00
line 002: 0x474F3020 : 0x00 0x00
line 003: 0x474F3030 : 0x00 0x00
line 004: 0x474F3040 : 0x00 0x00
line 005: 0x474F3050 : 0x00 0x00
line 006: 0x474F3060 : 0x00 0x00
line 007: 0x474F3070 : 0x00 0x00
line 008: 0x474F3080 : 0x00 0x00
line 009: 0x474F3090 : 0x00 0x00
line 010: 0x474F30A0 : 0x00 0x00
line 011: 0x474F30B0 : 0x00 0x00
line 012: 0x474F30C0 : 0x00 0x00
line 013: 0x474F30D0 : 0x00 0x00
line 014: 0x474F30E0 : 0x00 0x00
```

Finally ,Give the data to the program

2. Implementation

Software Model: Helping the programmer

Step 1: initialization:

The programmer defines a memory area as “randomized access” to protect sensitive data.



2. Implementation

Software Model: Helping the programmer

Step 1: initialization:

Note :

After the initialization- the **programmer** will access the memory as usual and **is oblivious** to the underlying permutation and actual “real” addresses.



2. Implementation

Behind the scenes:

A spin on a well-known meme

If the programmer has data in address **0xXXXXXX** he can expect to read back the same data from **0xXXXXXX**

Programmer's View



Actual Scrambled Data In RAM



2. Implementation

Memory access flow:

► Current R –mapping rule

History table

Line index	Original set
0	0x6E
1	0x12
2	0x19
3	0x28
4	0xFD
5	0x41
6	0x56

Memory area:

Memory arr:

```
line 000: 0x474F3000 : 0x00 0x00 0x00 0x00  
line 001: 0x474F3010 : 0x00 0x00 0x00 0x00  
line 002: 0x474F3020 : 0x00 0x00 0x00 0x00  
line 003: 0x474F3030 : 0x00 0x00 0x00 0x00  
line 004: 0x474F3040 : 0x00 0x00 0x00 0x00  
line 005: 0x474F3050 : 0x00 0x00 0x00 0x00  
line 006: 0x474F3060 : 0x00 0x00 0x00 0x00  
line 007: 0x474F3070 : 0x00 0x00 0x00 0x00  
line 008: 0x474F3080 : 0x00 0x00 0x00 0x00  
line 009: 0x474F3090 : 0x00 0x00 0x00 0x00  
line 010: 0x474F30A0 : 0x00 0x00 0x00 0x00  
line 011: 0x474F30B0 : 0x00 0x00 0x00 0x00  
line 012: 0x474F30C0 : 0x00 0x00 0x00 0x00  
line 013: 0x474F30D0 : 0x00 0x00 0x00 0x00  
line 014: 0x474F30E0 : 0x00 0x00 0x00 0x00  
* * *
```

2. Implementation

Memory access flow:

Programmer access an address within the array : 0xTTTTTSSO

2. Permutation on set bits – to randomize the cache line access

2. New address is : 0xTTTTTZZO

2. Implementation

Memory access flow:

2. New address is : 0xTTTTTZZO

3. Calculate line index in memory area

$$\text{line index} = \frac{\text{New Address} - \text{base address}}{2^{\text{size of offset field in bits}}}$$

3.b. Check history table : Is the data we are looking for stored at the
“New Address” line?

Is History table[line index] == original set ?

No!

Yes!

Next Slide

6. DONE!
Read the data at new address –
return the data to the program

2. Implementation

Memory access flow: Fixing the mapping

3.b. Check history table : Is the data we are looking for stored at the “**New Address**” line?

NO!



4. Search history table to find the line in which the data is (linearly or fast search)



5. Fix Mapping- swap the memory lines and update history table



6. DONE!

Read the data at new address –
return the data to the program

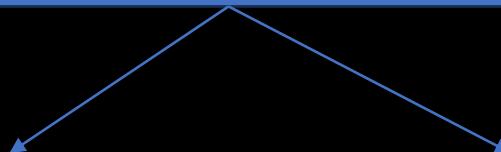
2. Implementation

Testing correctness and injectivity in C

Test clearBits ,Cswap2bit, Cswap8bit,

Test permute address for injectivity : $2^8 * 2^{20}$

Test reading back memory from array : $2^8 * 2^{20} * 2^{2^{12}}$



Without permutation (R=0)

With permutation (20 bit R)



2. Implementation

Tie HDL acceleration

3 Tie operation

8 Tie functions

C_swap(conditional swap)



Line calculation

Address calculation

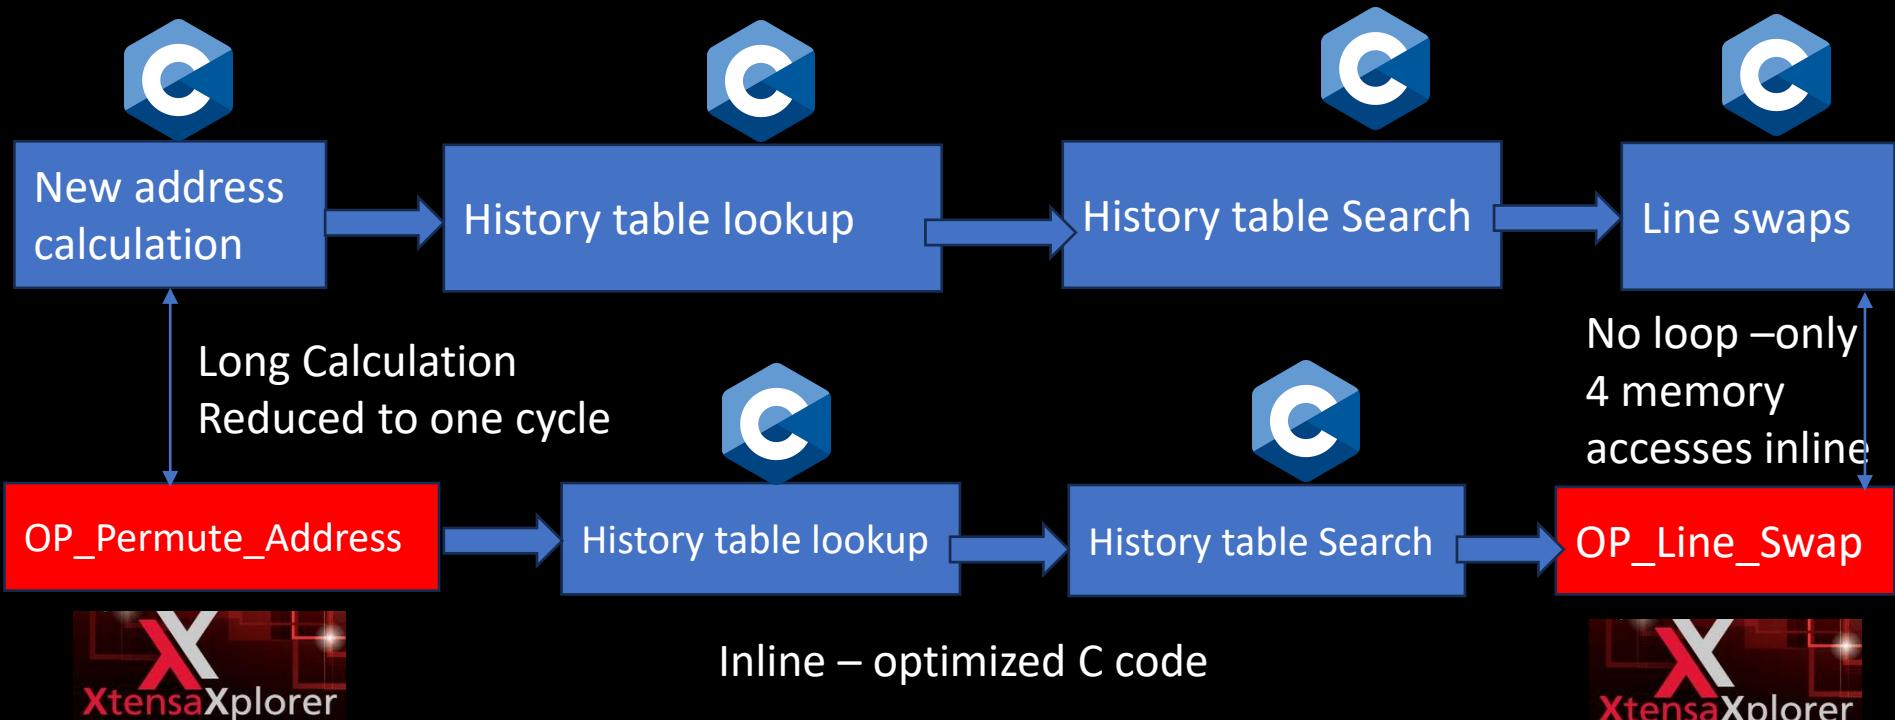
Vectorized 128 bit swap



Brute force testing :Comparison with C functions to verify correctness

2. Implementation

Tie HDL acceleration Memory access flow speedup



Note: line index calculation, history swap and more were optimized by analyzing disassembly with Oren

3. Statistical Analysis



3. Stastical Analysis

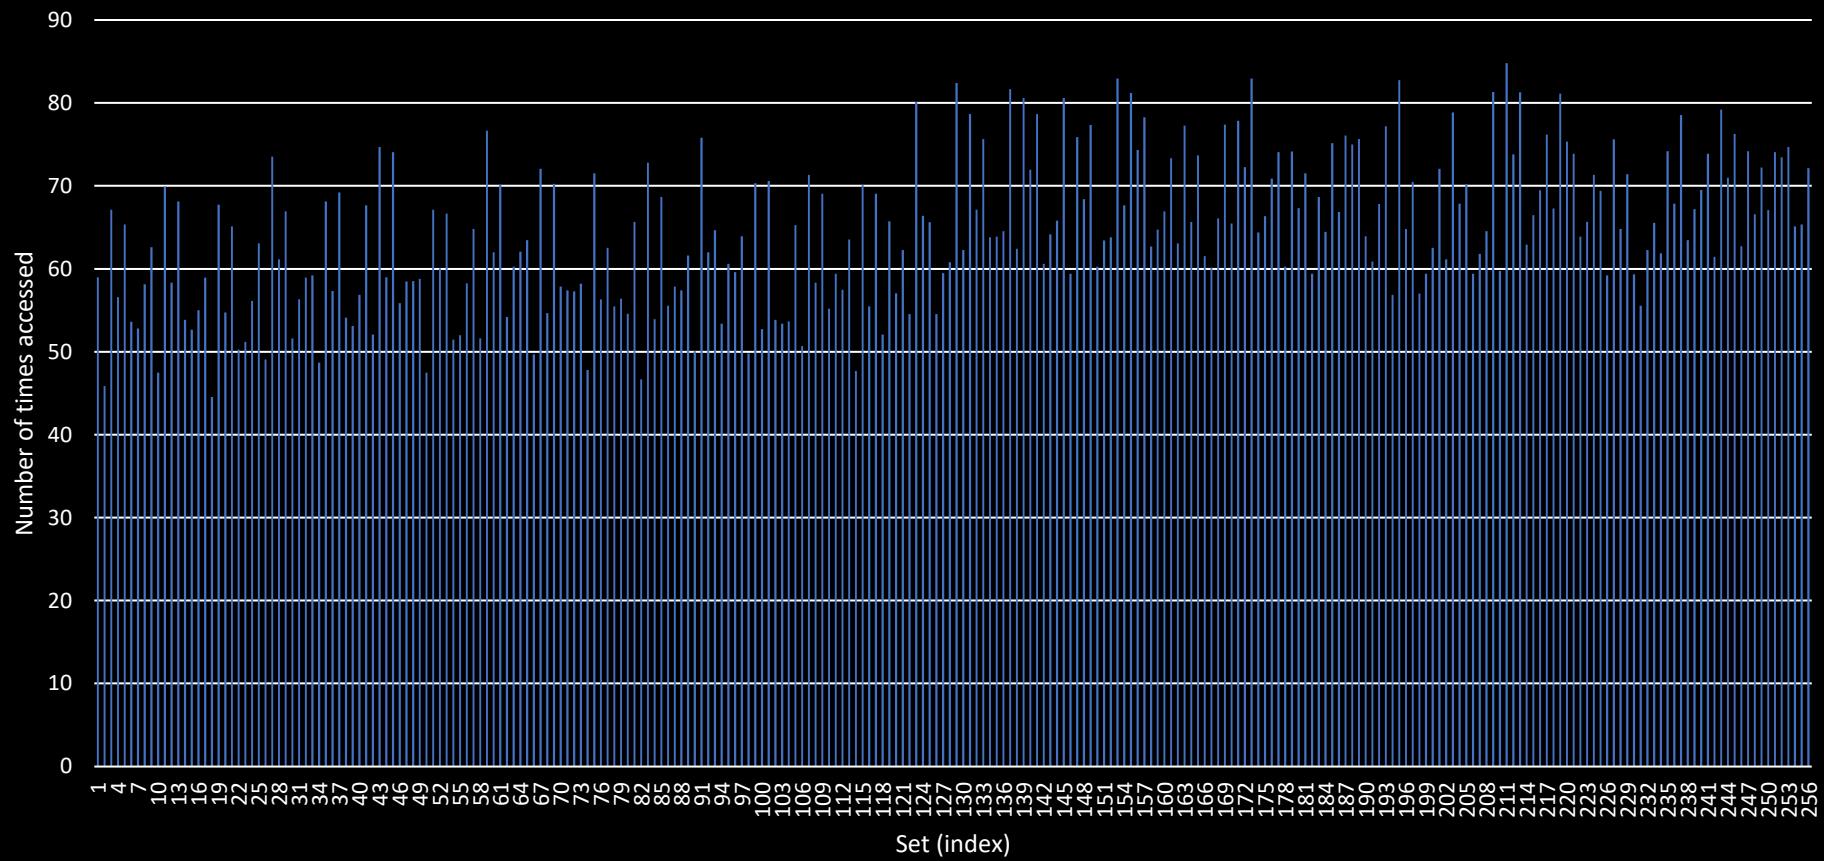
We want to prove that :

1. Uniform distribution (cryptographic security) – attacker cannot learn from set accesses
2. Show set access pattern over time – see how it masks the “real” set accesses

3. Statistical Analysis

- First test : Running a 15,000 sequential reads and counting how many times each set was accessed:

Average number of accesses per each set when r replacements =10



3. Statistical Analysis

- First test : Running a N=15,000 sequential reads and counting how many times each set was accessed:

- We used statistical tests to verify that the expected value for the average number of accesses per set is about $n * 1/256$.
- Coefficient of variation : of differences from mean
- As well as Pearson's Chi-squared test to check the proportions are matching the expected 1/256 proportion.

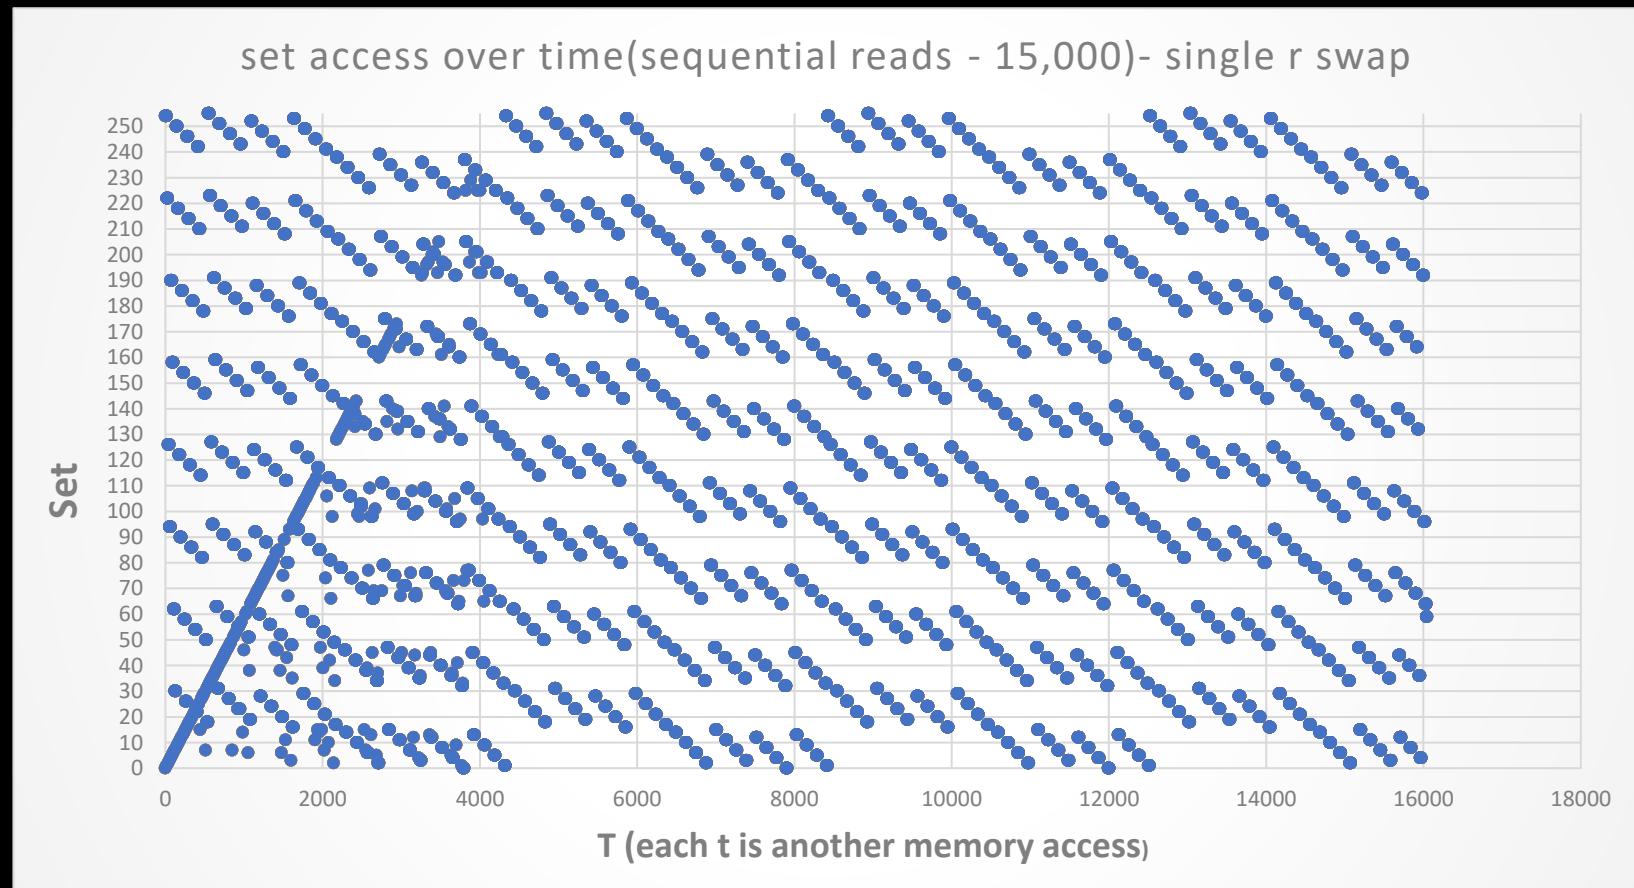


Set index sequential	Set index number of accesses	Experiment/Sample index out of N					
		1	2	3	4	5	6
0	81	86	38	40	52	67	
1	51	69	24	23	38	38	
2	81	69	38	68	52	66	
3	67	68	22	67	38	68	
4	81	69	38	38	52	69	
5	52	57	24	23	42	52	
6	65	70	12	52	54	40	
7	53	69	20	57	53	53	



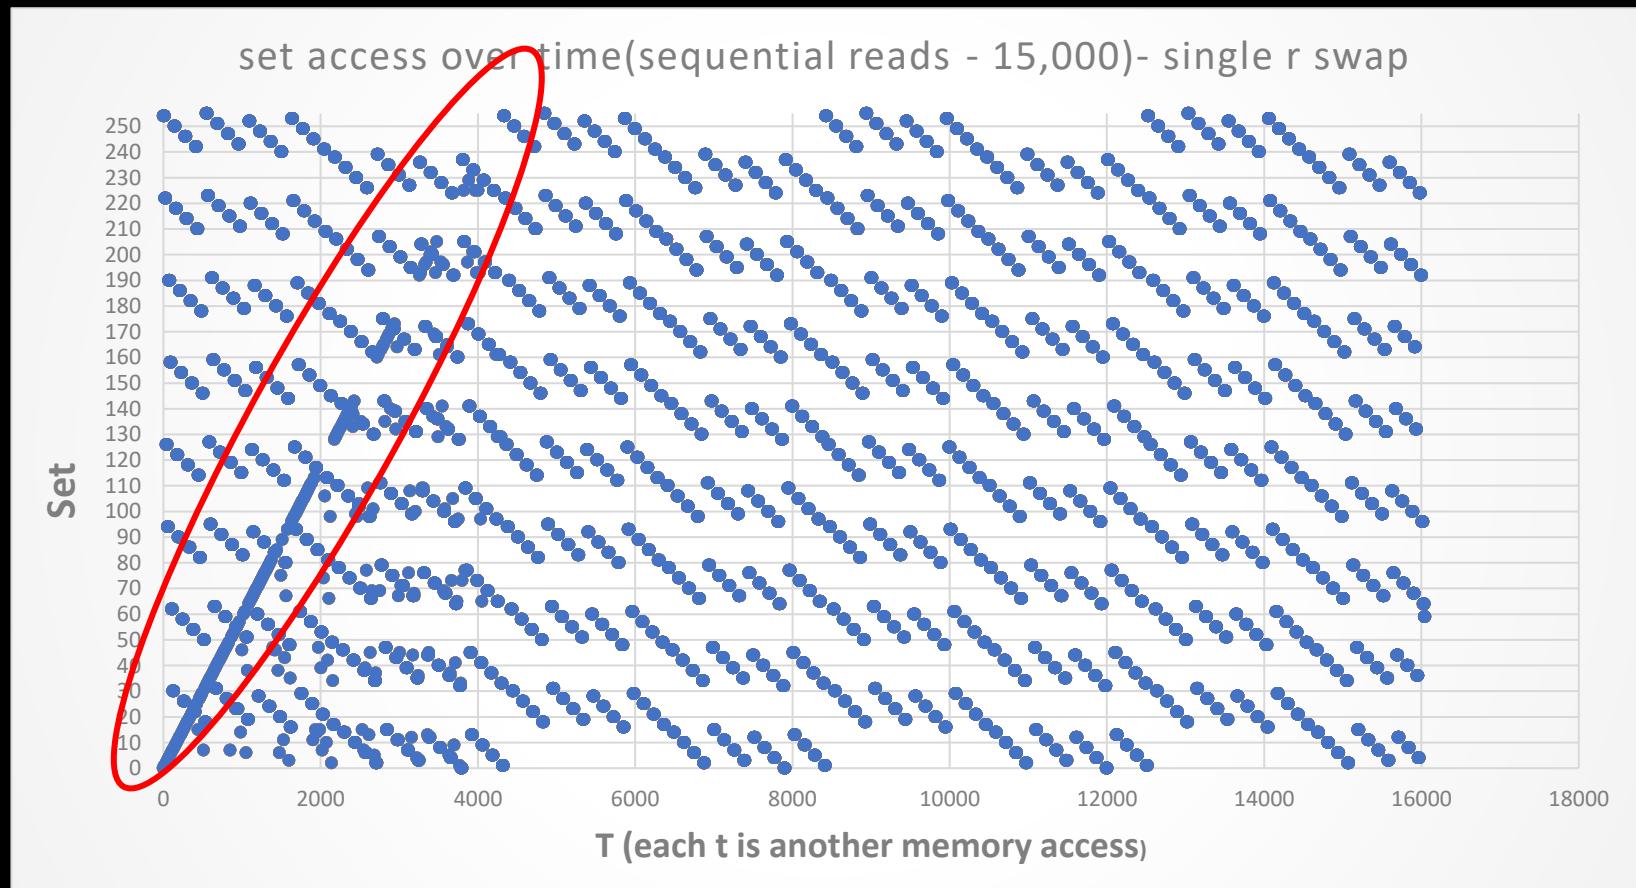
3. Statistical Analysis

- Second test – How frequently should we swap the permutation?



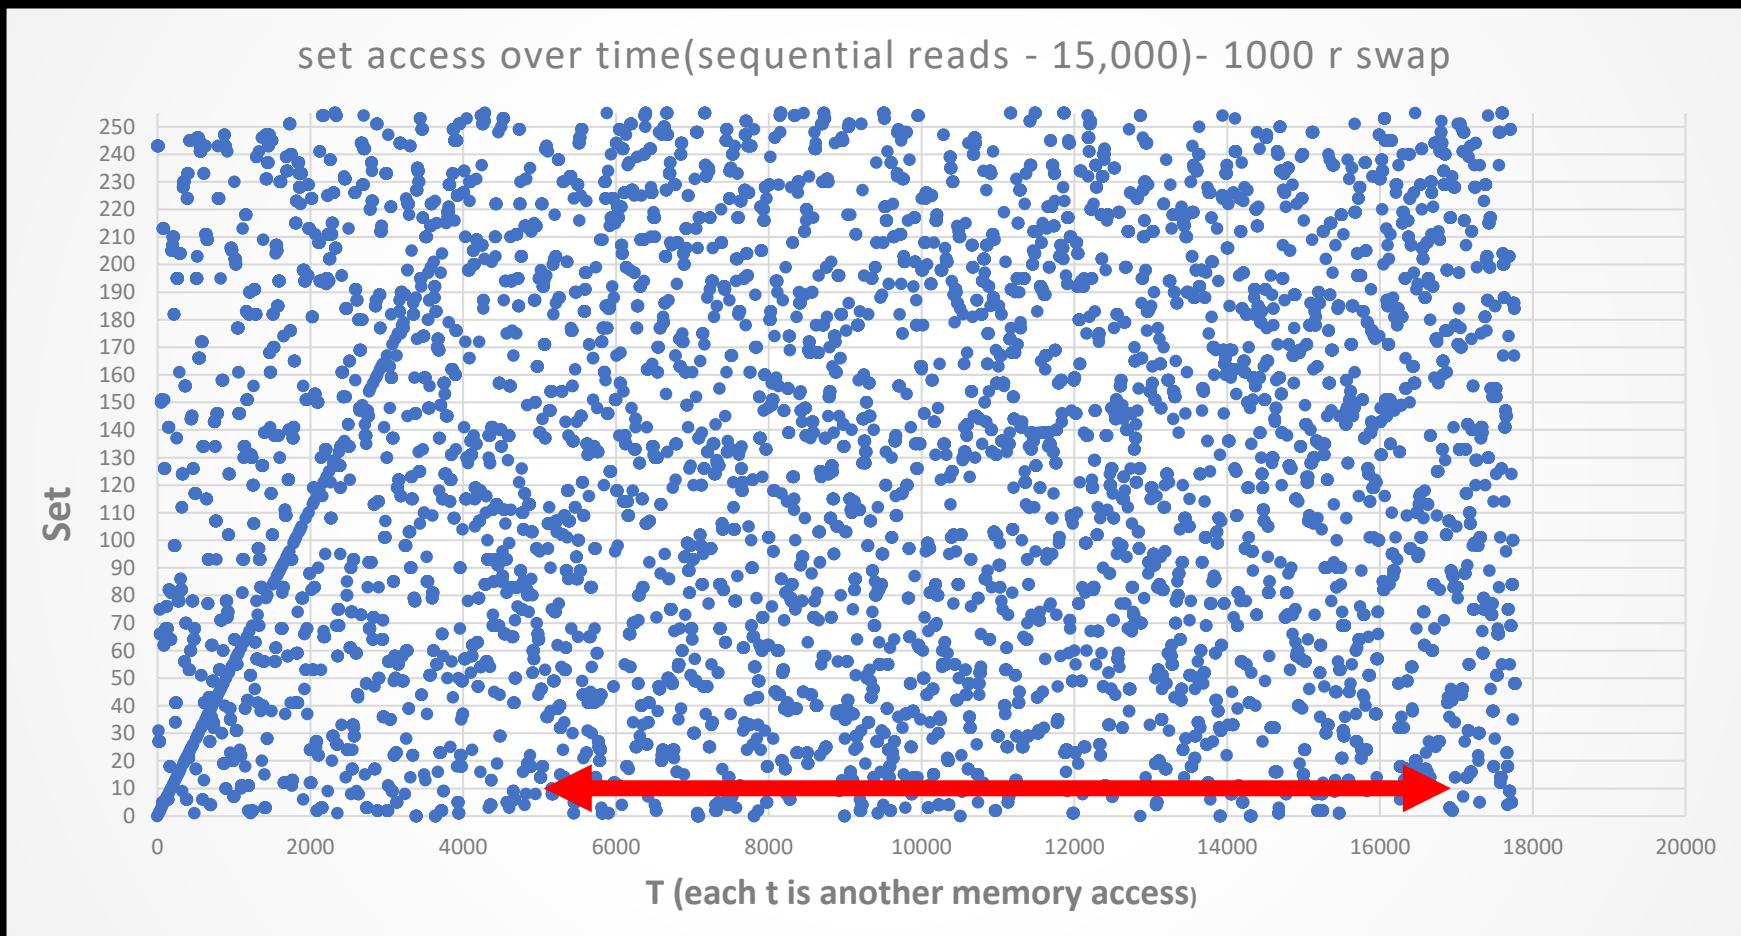
3. Statistical Analysis

- Second test – How frequently should we swap the permutation?
**Pay attention to the linear line in the beginning:
Marks the initial shuffling of original memory mapping.**



3. Statistical Analysis

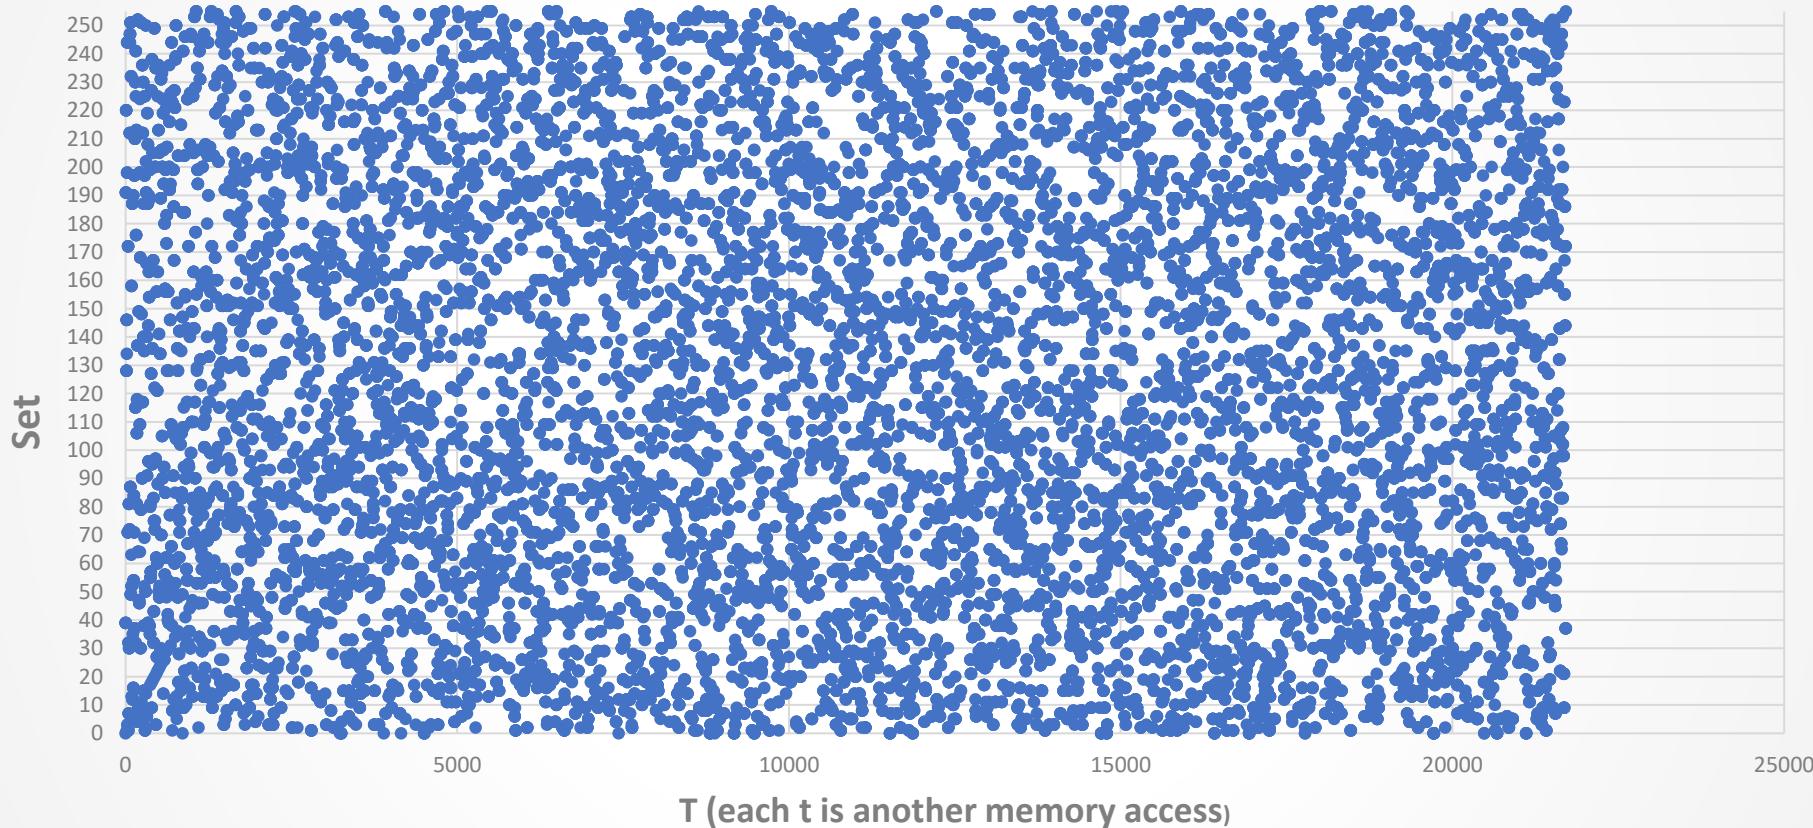
- Second test – How frequently should we swap the permutation?
If we increase the frequency of changing R – the access pattern is not visible anymore



3. Statistical Analysis

- Second test – How frequently should we swap the permutation?
If we increase the frequency of changing R – the access pattern is not visible anymore

set access over time(sequential reads - 15,000)- 5000 r swap



3. Statistical Analysis

- Third test –Protection performance penalty:

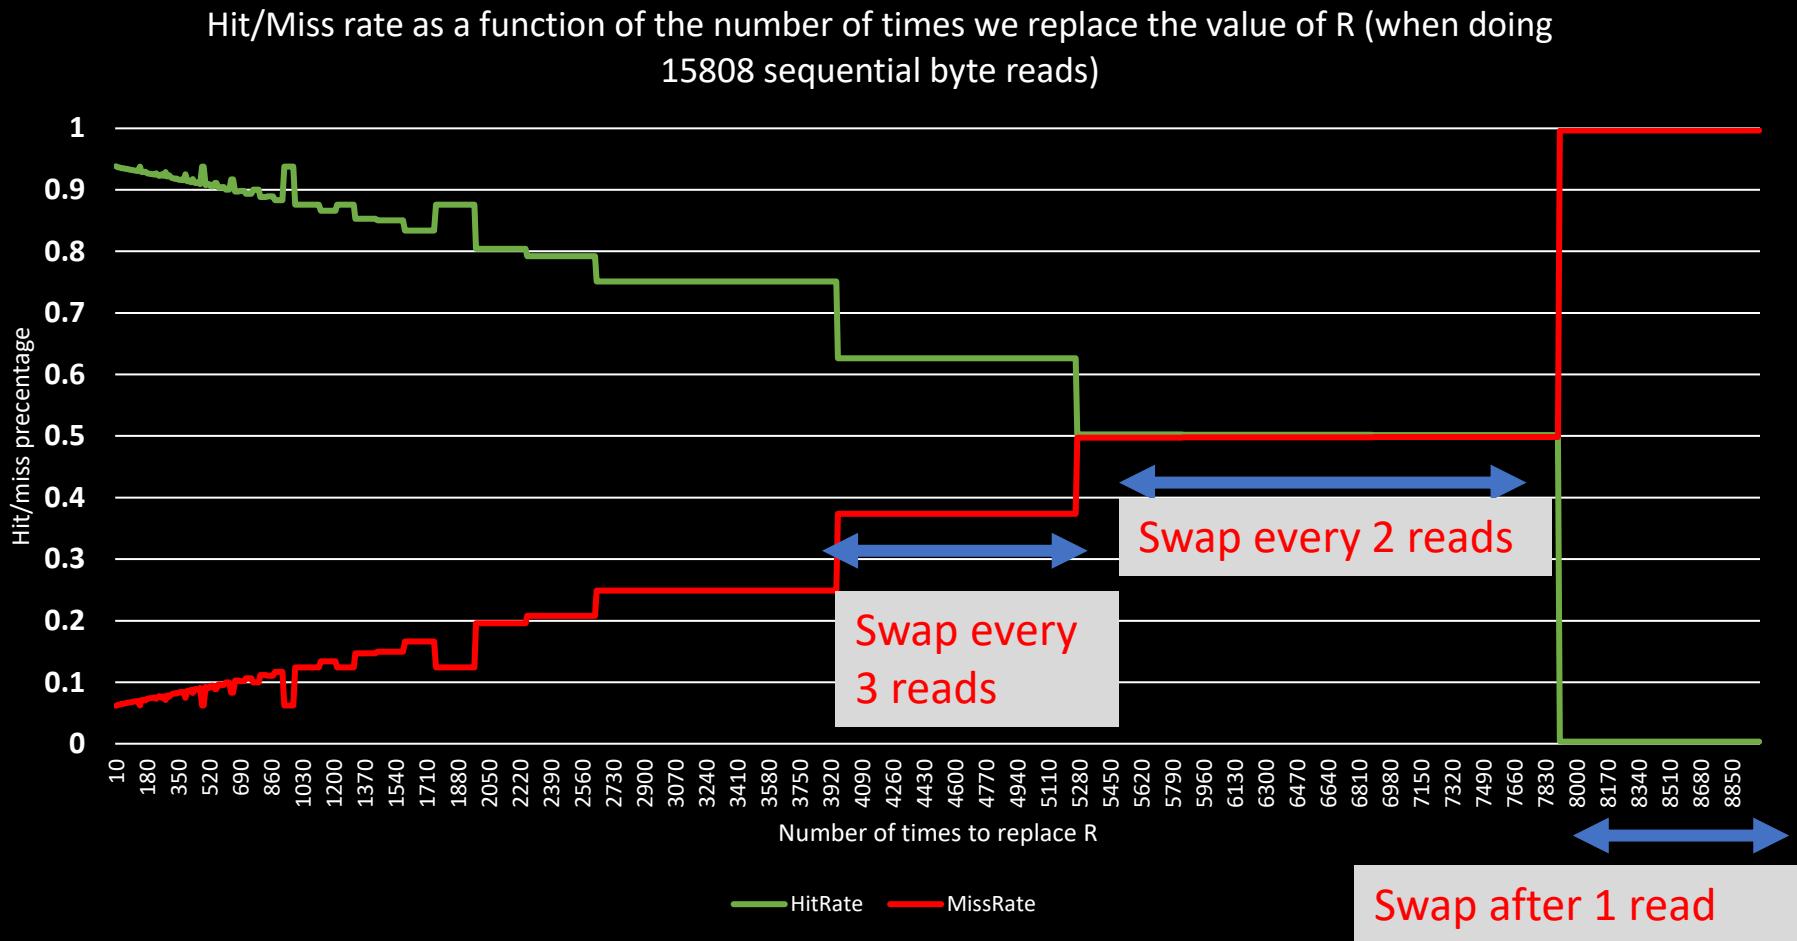
How many memory swaps will happen due to “misses” in history table-> when we increase the frequency of permutation swaps ?
(N=15,000 sequential reads)



3. Statistical Analysis

- Third test –Protection performance penalty:

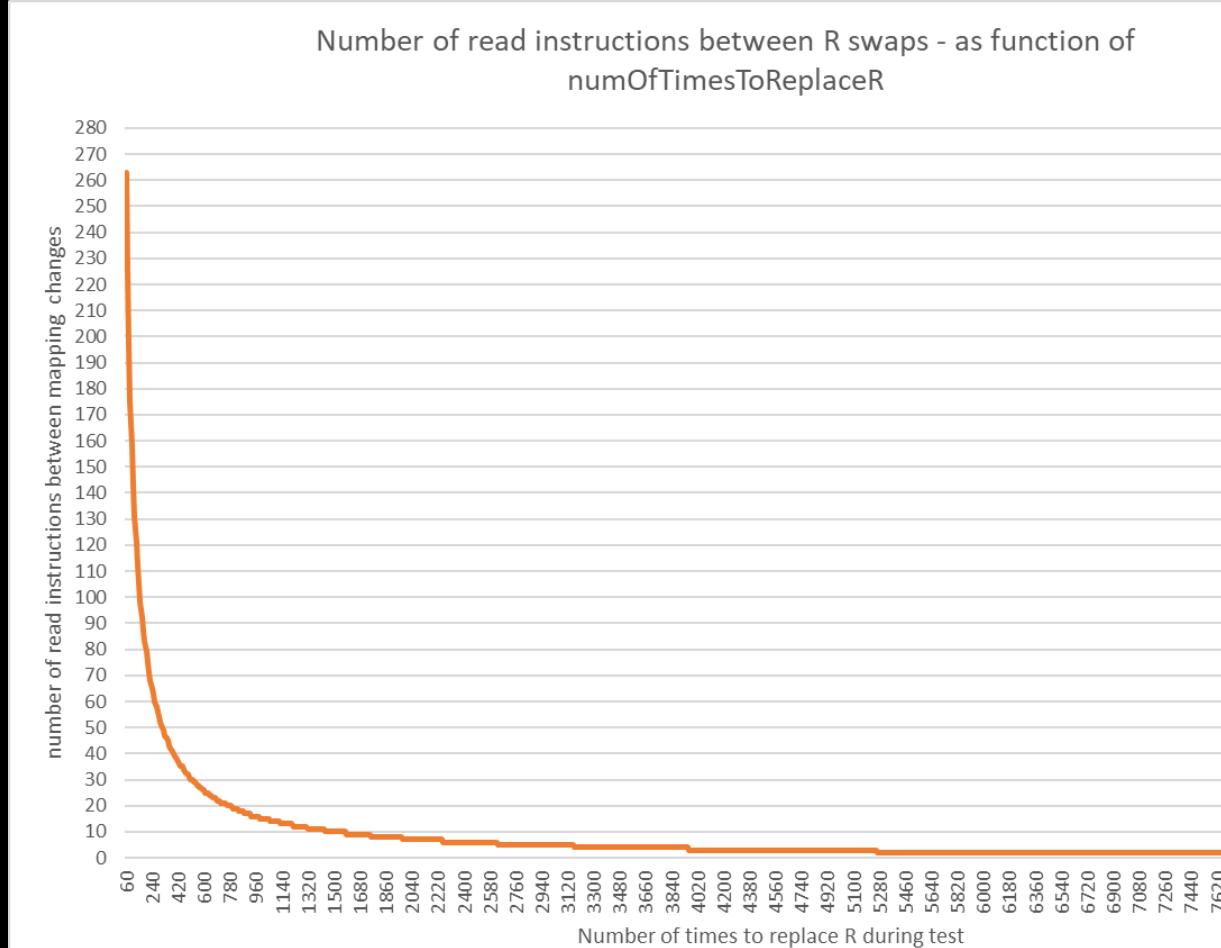
How many memory swaps will happen due to “misses” in history table-> when we increase the frequency of permutation swaps ?
(N=15,000 sequential reads)



3. Statistical Analysis

- Reaching a conclusion – how often should we swap the permutation for address mapping?

If We compare the Number of read instructions between R swaps - as function of numOfTimesToReplaceR .
We can answer the question : what is the optimal frequency range to swap the permutation (swapping R)?



3. Statistical Analysis

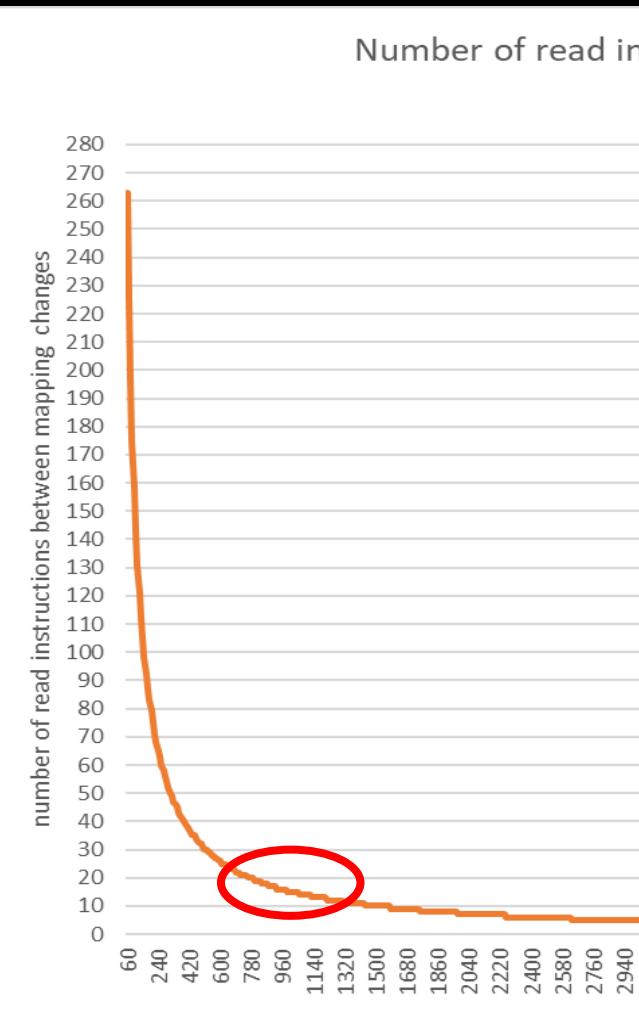
- Reaching a conclusion – how often should we swap the permutation for address mapping?



Conclusion :
In the range of
900-1000

We get over
90% 'Hit'

Which is an R-
value swap
every 15-17
memory
accesses.

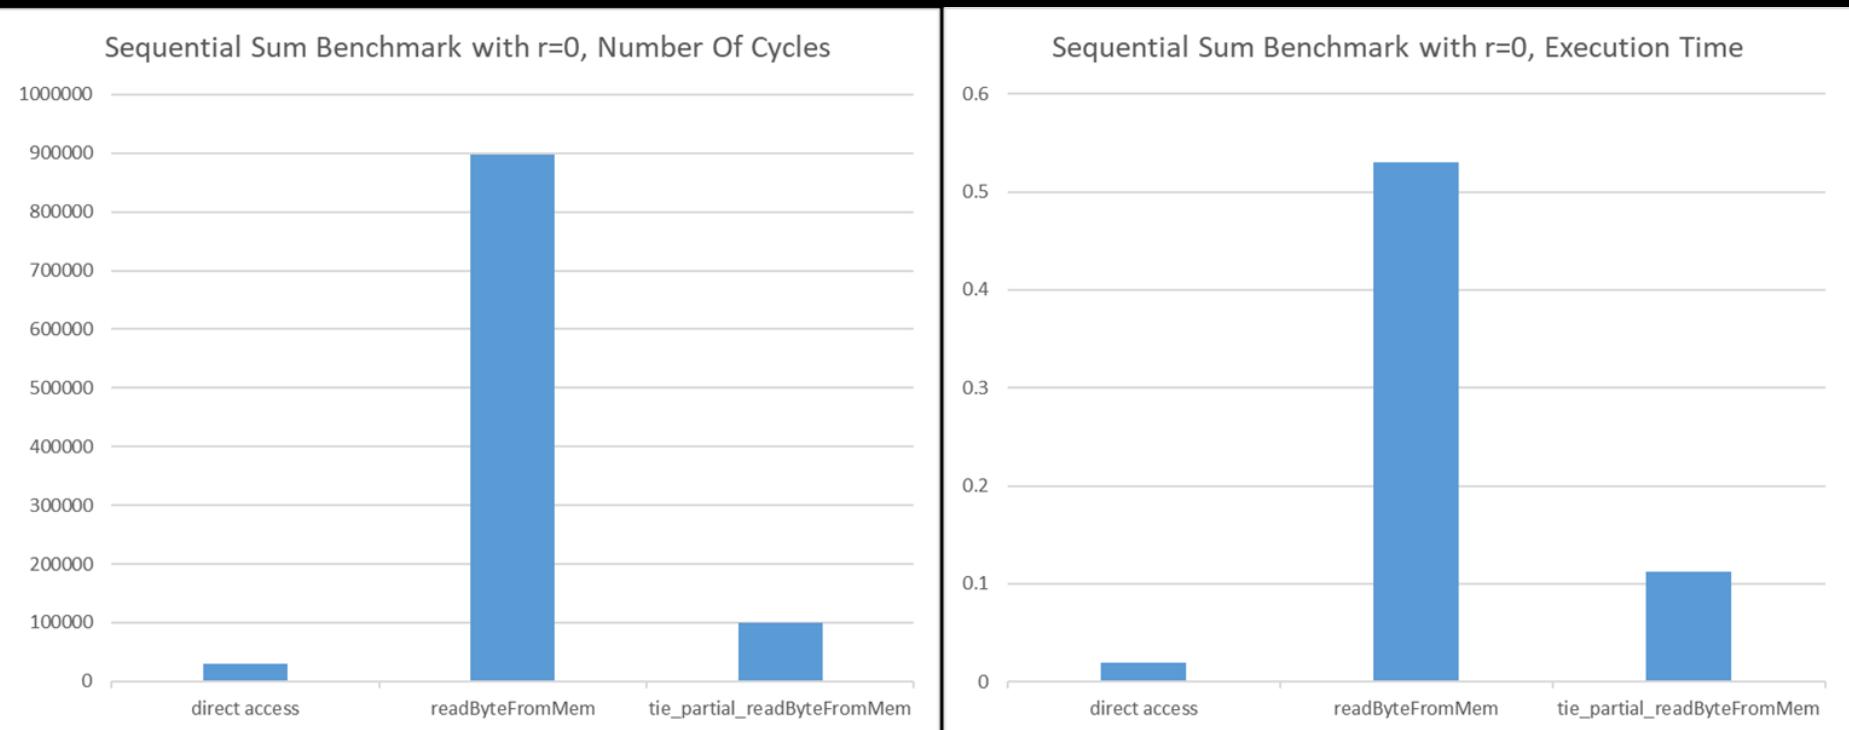


4. Results and Benchmarks



4. Benchmarks

1. Sequential Sum
Benchmark: accessing
whole memory
(4096 memory accesses)

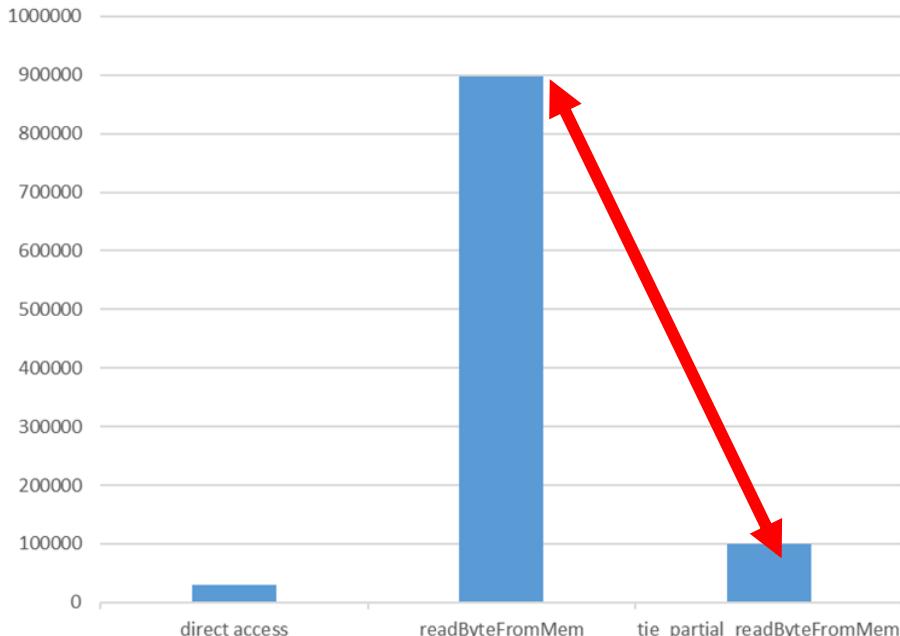


4. Benchmarks

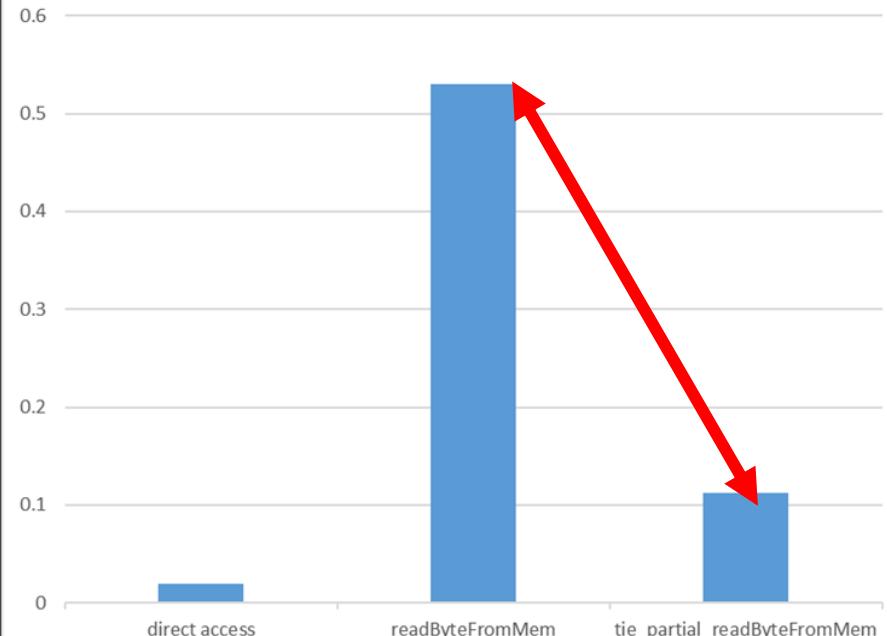
1. Sequential Sum
Benchmark: accessing
whole memory
(4096 memory accesses)

Significant acceleration ! Tie operations saves us :

Sequential Sum Benchmark with r=0, Number Of Cycles



Sequential Sum Benchmark with r=0, Execution Time



4. Benchmarks

1. Sequential Sum
Benchmark: accessing whole memory
(4096 memory accesses)

Significant acceleration ! Tie operations saves us :



Cycles :

If $R=0 \rightarrow$ Reduced by 89%

If $R \neq 0 \rightarrow$ On average :
Reduced by 60%

Due to history table
searches and swaps

Time :

If $R=0 \rightarrow$ speedup of 4.7 times faster!

If $R \neq 0 \rightarrow$ average speedup of 2.3 !



4. Benchmarks

2. Binary Search Benchmark-Results

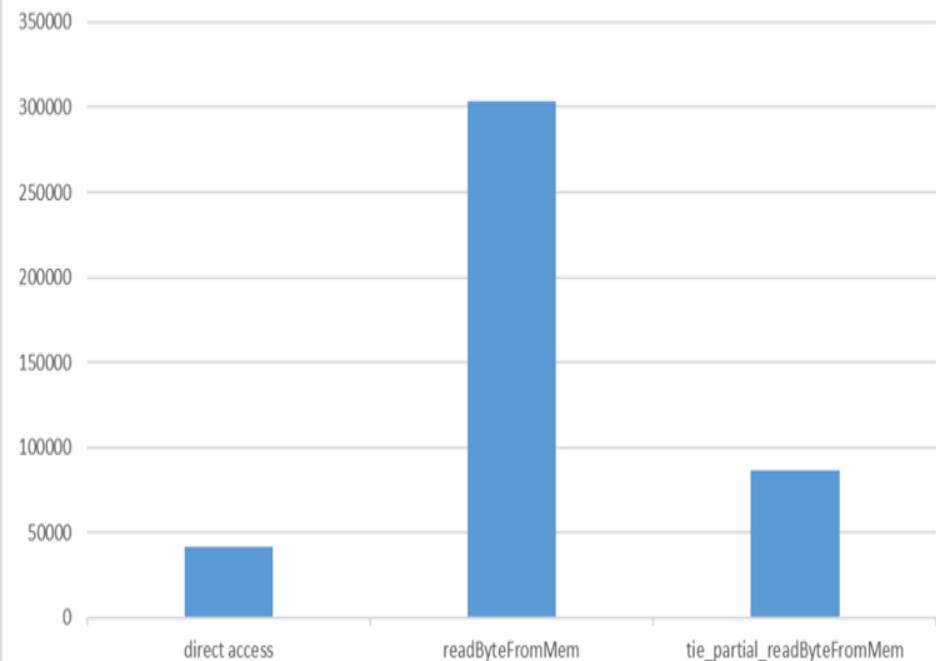
Non-sequential reads **benchmark** :
unknown access pattern

Binary Search										
Search 23	0	1	2	3	4	5	6	7	8	9
	2	5	8	12	16	23	38	56	72	91
23 > 16 take 2 nd half	L=0	1	2	3	M=4	5	6	7	8	H=9
	2	5	8	12	16	23	38	56	72	91
23 < 56 take 1 st half	0	1	2	3	4	L=5	6	M=7	8	H=9
	2	5	8	12	16	23	38	56	72	91
Found 23, Return 5	0	1	2	3	4	L=5, M=5	H=6	7	8	9
	2	5	8	12	16	23	38	56	72	91

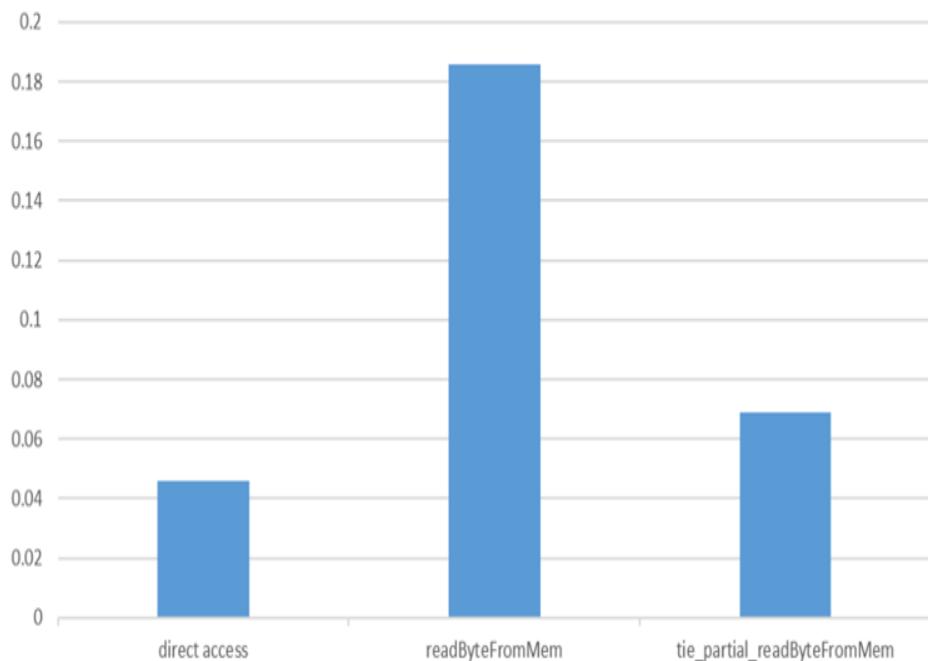
4. Benchmarks

2. Binary Search Benchmark-Results

Binary Search Benchmark, average of different values and R values
(100 iterations of binary search with same values)



Binary Search Benchmark, average of different values and R values
(100 iterations of binary search with same values)



4. Benchmarks

2. Binary Search Benchmark-Results

Significant acceleration ! Tie operations saves us :



Cycles :

If $R \neq 0$ -> On average :
Reduced by 72%



Time :

If $R \neq 0$ -> average
speedup of 2.7 !

Time Penalty

- Disadvantages, downsides and cost of using our information leakage protection:
- If a programmer deploys our solution :the average runtime for sequential 4096 reads will increase by up to 14 times (on average).
- If a programmer deploys our solution – the average runtime for Binary search **will increase by 1.5 times**. The cycle count will be 2.1 times higher



Conclusion

- The system works
- It protects from leakage by making “SET” cache line accesses randomized ($0 = 1/256$).
- TIE HDL gives a significant acceleration to our algorithm
- Programmers should choose wisely when to implement this solution – only for critical safe/cryptographic data – to avoid the performance penalty.



שאלות?

תולדת רباح